

predictive-maintenance-1

February 24, 2024

```
[1]: #Title: Predictive Maintenance for Industrial Equipment Using Machine Learning
```

```
[1]: import pandas as pd
import numpy as np
```

```
[2]: data = pd.read_csv('Dataset - Sheet1.csv')
```

```
[3]: data
```

```
[3]:
```

| | Temperature | Pressure | Vibration | Humidity | Failure |
|------|-------------|-----------|-----------|----------|---------|
| 0 | 23.75594 | 111.47800 | 0.590529 | 42.99730 | 0 |
| 1 | 18.38782 | 108.31190 | 0.561806 | 26.16201 | 0 |
| 2 | 25.86647 | 86.78422 | 0.434486 | 40.09598 | 0 |
| 3 | 21.07013 | 104.74830 | 0.602397 | 38.64469 | 0 |
| 4 | 22.74969 | 95.66371 | 0.553688 | 42.28409 | 0 |
| ... | ... | ... | ... | ... | ... |
| 8732 | 21.98014 | 111.96370 | 0.505857 | 66.63030 | 0 |
| 8733 | 28.56523 | 106.07030 | 0.379595 | 57.72110 | 0 |
| 8734 | 26.39877 | 90.06860 | 0.602586 | 51.97461 | 0 |
| 8735 | 34.30743 | 104.64770 | 0.430827 | 60.13049 | 0 |
| 8736 | 21.56679 | 97.80628 | 0.440026 | 29.18022 | 0 |

[8737 rows x 5 columns]

```
[4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8737 entries, 0 to 8736
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   Temperature     8736 non-null   float64
1   Pressure        8735 non-null   float64
2   Vibration       8736 non-null   float64
3   Humidity        8737 non-null   float64
4   Failure         8737 non-null   int64  
dtypes: float64(4), int64(1)
memory usage: 341.4 KB
```

```
[5]: data.isnull().sum()
```

```
[5]: Temperature    1
      Pressure      2
      Vibration     1
      Humidity      0
      Failure       0
      dtype: int64
```

```
[6]: data = data.drop_duplicates()
```

```
[7]: data
```

```
[7]:
```

| | Temperature | Pressure | Vibration | Humidity | Failure |
|------|-------------|-----------|-----------|----------|---------|
| 0 | 23.75594 | 111.47800 | 0.590529 | 42.99730 | 0 |
| 1 | 18.38782 | 108.31190 | 0.561806 | 26.16201 | 0 |
| 2 | 25.86647 | 86.78422 | 0.434486 | 40.09598 | 0 |
| 3 | 21.07013 | 104.74830 | 0.602397 | 38.64469 | 0 |
| 4 | 22.74969 | 95.66371 | 0.553688 | 42.28409 | 0 |
| ... | ... | ... | ... | ... | ... |
| 8732 | 21.98014 | 111.96370 | 0.505857 | 66.63030 | 0 |
| 8733 | 28.56523 | 106.07030 | 0.379595 | 57.72110 | 0 |
| 8734 | 26.39877 | 90.06860 | 0.602586 | 51.97461 | 0 |
| 8735 | 34.30743 | 104.64770 | 0.430827 | 60.13049 | 0 |
| 8736 | 21.56679 | 97.80628 | 0.440026 | 29.18022 | 0 |

[8737 rows x 5 columns]

```
[8]: data = data.fillna(data.mean()) #Null values replaced by mean
```

```
[9]: data.isnull().sum()
```

```
[9]: Temperature    0
      Pressure      0
      Vibration     0
      Humidity      0
      Failure       0
      dtype: int64
```

```
[10]: import pandas as pd
      from scipy.stats import zscore
      data_standardized = data.apply(zscore)
```

```
[11]: data_standardized # for data standardisation of data zscore applied
```

```
[11]:
```

| | Temperature | Pressure | Vibration | Humidity | Failure |
|---|-------------|----------|-----------|-----------|-----------|
| 0 | -0.241262 | 1.142597 | 0.923072 | -0.690062 | -0.226401 |

```

1      -1.316431  0.828578  0.634674 -2.360345 -0.226401
2       0.181451 -1.306573 -0.643702 -0.977911 -0.226401
3      -0.779197  0.475135  1.042234 -1.121898 -0.226401
4      -0.442801 -0.425890  0.553164 -0.760822 -0.226401
...
8732   -0.596933  1.190770  0.072910  1.654642 -0.226401
8733    0.721980  0.606253 -1.194843  0.770732 -0.226401
8734    0.288065 -0.980823  1.044132  0.200605 -0.226401
8735    1.872072  0.465157 -0.680441  1.009776 -0.226401
8736   -0.679722 -0.213387 -0.588077 -2.060898 -0.226401

```

[8737 rows x 5 columns]

```
[12]: data_standardized.columns
```

```
[12]: Index(['Temperature', 'Pressure', 'Vibration', 'Humidity', 'Failure'],
dtype='object')
```

```
[13]: import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
```

```
[14]: #Correlation Analysis (Numeric-Numeric)
data_standardized.corr()
```

```
[14]:
```

| | Temperature | Pressure | Vibration | Humidity | Failure |
|-------------|-------------|-----------|-----------|-----------|-----------|
| Temperature | 1.000000 | -0.003997 | -0.003345 | 0.018575 | -0.005177 |
| Pressure | -0.003997 | 1.000000 | -0.011699 | -0.003345 | -0.010030 |
| Vibration | -0.003345 | -0.011699 | 1.000000 | -0.011047 | -0.006912 |
| Humidity | 0.018575 | -0.003345 | -0.011047 | 1.000000 | -0.004122 |
| Failure | -0.005177 | -0.010030 | -0.006912 | -0.004122 | 1.000000 |

```
[15]: X = data_standardized.drop("Failure", axis = 1).values
```

```
[16]: X
```

```
[16]: array([[ -0.2412619 ,  1.14259744,  0.92307166, -0.6900623 ],
        [ -1.31643057,  0.82857835,  0.63467401, -2.36034454],
        [  0.1814514 , -1.30657295, -0.64370183, -0.97791141],
        ...,
        [  0.28806457, -0.98082265,  1.04413181,  0.20060496],
        [  1.87207228,  0.46515686, -0.68044058,  1.00977555],
        [ -0.67972183, -0.21338661, -0.58807662, -2.0608984 ]])
```

```
[17]: y = data_standardized['Failure'].values
```

```
[18]: y
```

```
[18]: array([-0.2264009, -0.2264009, -0.2264009, ..., -0.2264009, -0.2264009,
          -0.2264009])
```

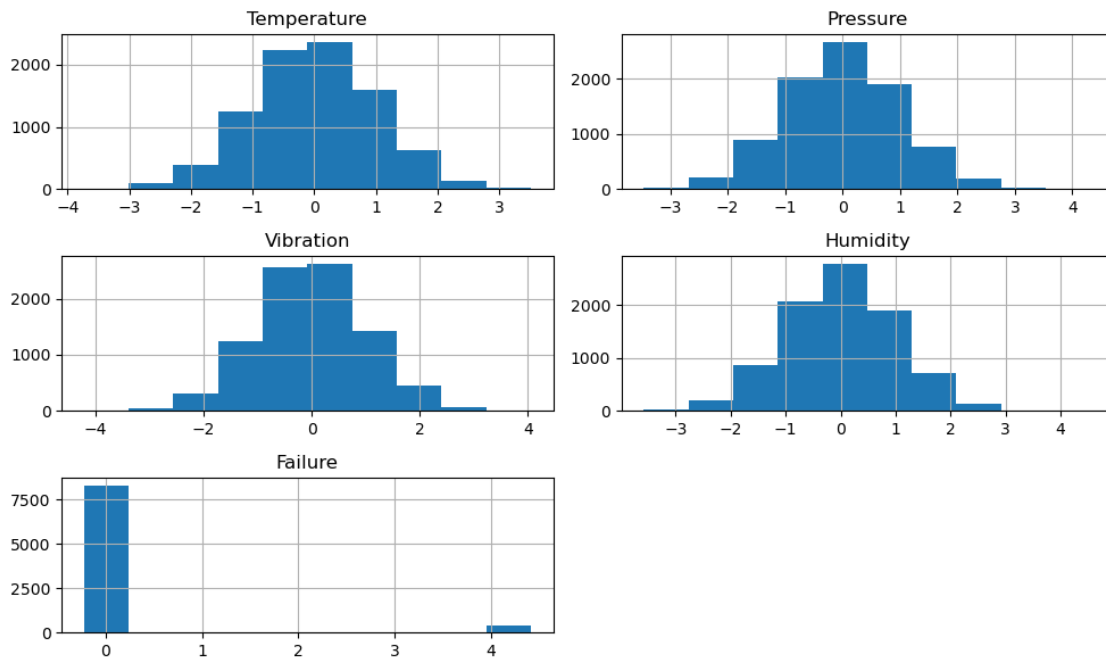
```
[19]: data_standardized
```

```
[19]:      Temperature  Pressure  Vibration  Humidity  Failure
0      -0.241262   1.142597   0.923072 -0.690062 -0.226401
1      -1.316431   0.828578   0.634674 -2.360345 -0.226401
2       0.181451  -1.306573  -0.643702 -0.977911 -0.226401
3      -0.779197   0.475135   1.042234 -1.121898 -0.226401
4      -0.442801  -0.425890   0.553164 -0.760822 -0.226401
...
8732    -0.596933   1.190770   0.072910  1.654642 -0.226401
8733     0.721980   0.606253  -1.194843  0.770732 -0.226401
8734     0.288065  -0.980823   1.044132  0.200605 -0.226401
8735     1.872072   0.465157  -0.680441  1.009776 -0.226401
8736    -0.679722  -0.213387  -0.588077 -2.060898 -0.226401
```

[8737 rows x 5 columns]


```
[20]: #for checking distribution of data
data_standardized.hist(figsize = (10,6))
plt.tight_layout()
plt.plot()
```

```
[20]: []
```

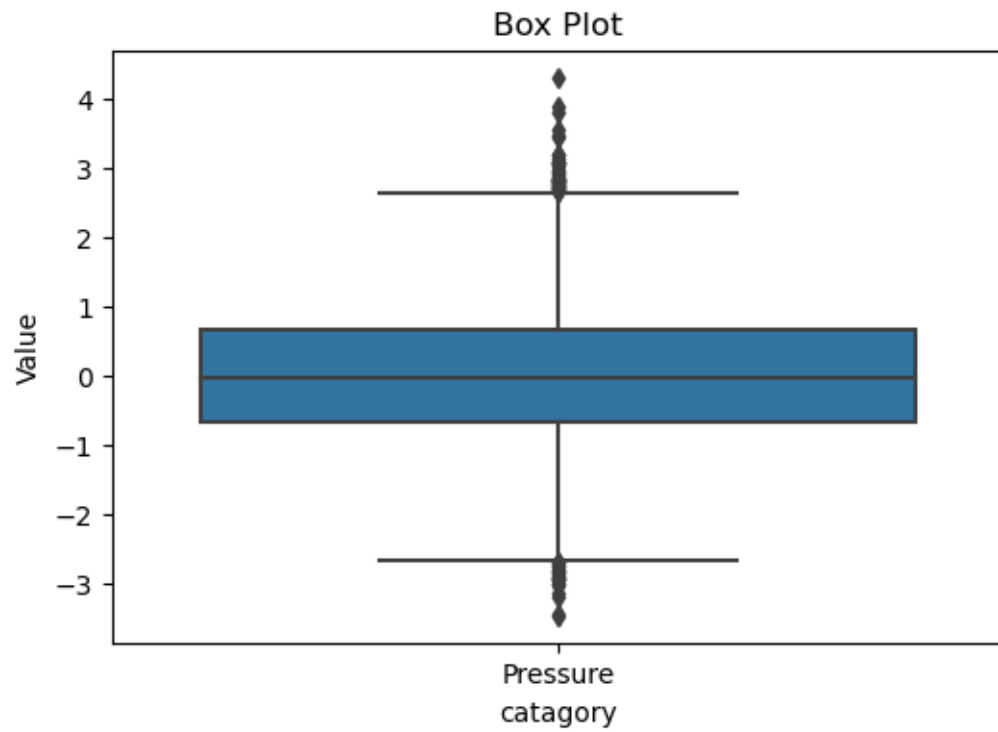


```
[21]: #outlier detection
import seaborn as sns
```

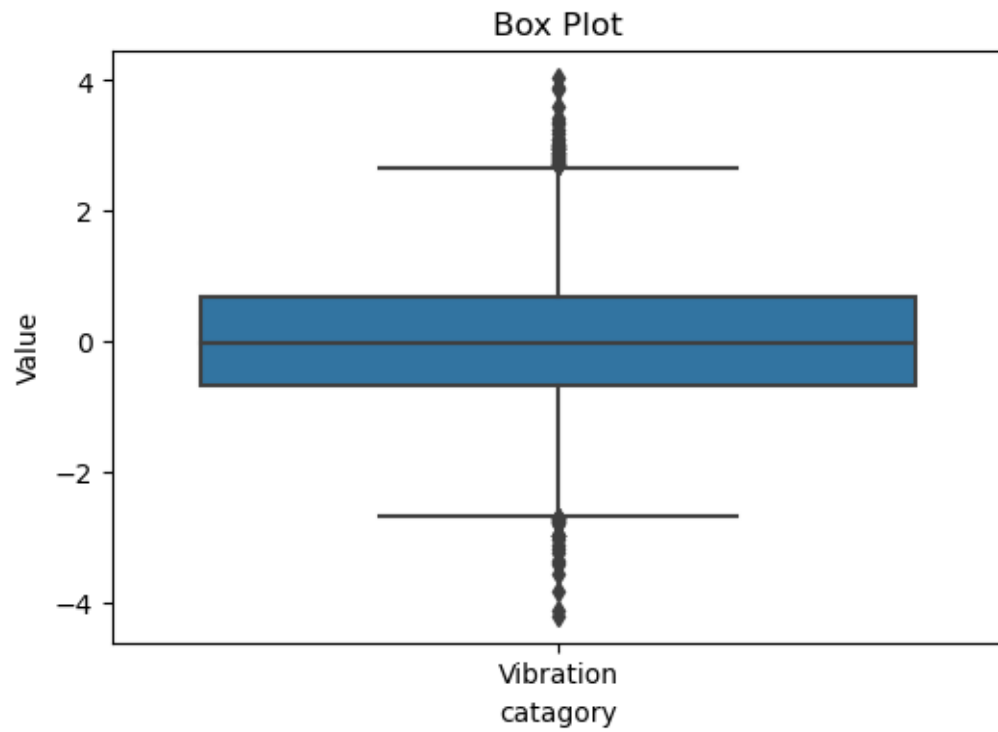
```
[22]: column_name = ['Temperature']
plt.figure(figsize = (6,4))
sns.boxplot(data = data_standardized[column_name])
plt.xlabel('catagory')
plt.ylabel('Value')
plt.title('Box Plot')
plt.show()
```



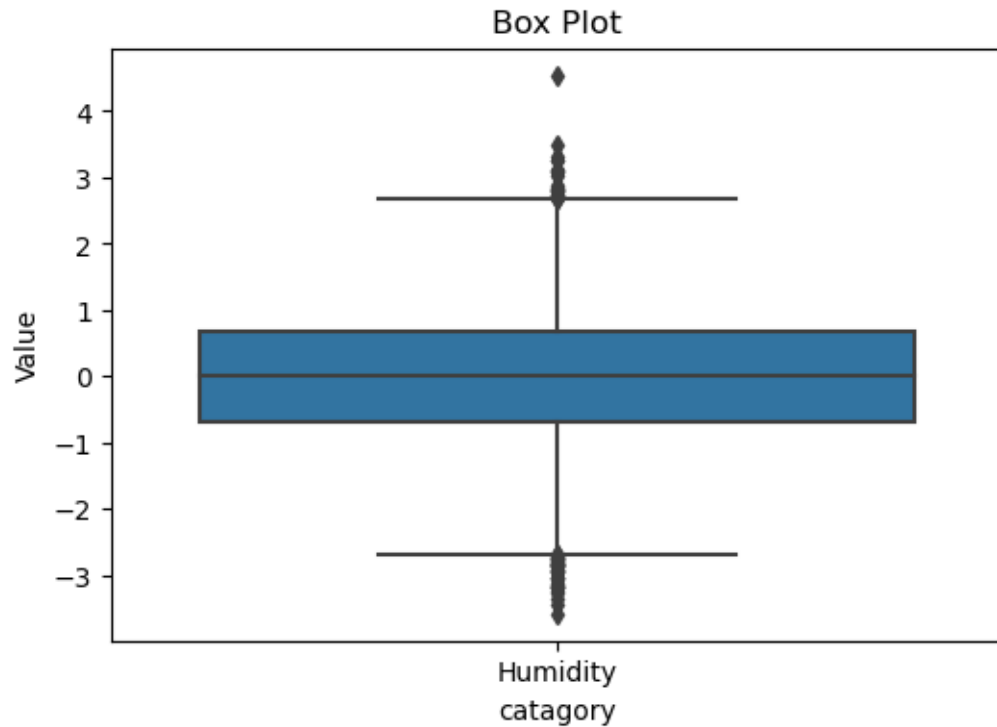
```
[23]: column_name = ['Pressure']
plt.figure(figsize = (6,4))
sns.boxplot(data = data_standardized[column_name])
plt.xlabel('catagory')
plt.ylabel('Value')
plt.title('Box Plot')
plt.show()
```



```
[24]: column_name = ['Vibration']  
plt.figure(figsize = (6,4))  
sns.boxplot(data = data_standardized[column_name])  
plt.xlabel('catagory')  
plt.ylabel('Value')  
plt.title('Box Plot')  
plt.show()
```



```
[25]: column_name = ['Humidity']  
plt.figure(figsize = (6,4))  
sns.boxplot(data = data_standardized[column_name])  
plt.xlabel('catagory')  
plt.ylabel('Value')  
plt.title('Box Plot')  
plt.show()
```



```
[26]: # Random Forest, Gradient Boosting, Support Vector Machines, etc.
```

```
[27]: #Random Forest
```

```
[83]: from sklearn.model_selection import train_test_split
```

```
[84]: X_train,X_test,y_train,y_test = train_test_split(X,y, test_size=0.2,
↳random_state=50)
```

```
[85]: X_train.shape
```

```
[85]: (6989, 4)
```

```
[86]: X_test.shape
```

```
[86]: (1748, 4)
```

```
[96]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score

# Split the dataset into training and testing sets
```



```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)

# Initialize the Random Forest Classifier
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)

# Train the classifier on the training data
rf_classifier.fit(X_train, y_train)

# Make predictions on the testing data
y_pred = rf_classifier.predict(X_test)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)

print("Accuracy:", accuracy)

```

Accuracy: 0.9490846681922197

```
[ ]: # Support Vector Machines
```

```
[56]: y = data['Failure'].values # checking failure states
```

```
[57]: y
```

```
[57]: array([0, 0, 0, ..., 0, 0, 0], dtype=int64)
```

```
[58]: from sklearn.model_selection import train_test_split
```

```
[59]: x_train,x_test,y_train,y_test = train_test_split(X,y,test_size=0.2,
↳random_state=50)
```

```
[60]: x_train.shape
```

```
[60]: (6989, 4)
```

```
[61]: x_test.shape
```

```
[61]: (1748, 4)
```

```
[62]: from sklearn.svm import SVC
```

```
[63]: cls = SVC(kernel="rbf")
```

```
[64]: cls.fit(x_train,y_train)
```

```
[64]: SVC()
```

```

[65]: ypred = cls.predict(x_test)

[66]: from sklearn.metrics import accuracy_score

[67]: accuracy_score(ypred,y_test)

[67]: 0.9559496567505721

[68]: # 'kernel': ['linear', 'poly', 'rbf', 'sigmoid'], # Kernel type for
      ↪ hyperparameter tuning

[69]: cls = SVC(kernel="linear")

[70]: cls.fit(x_train,y_train)

[70]: SVC(kernel='linear')

[71]: ypred = cls.predict(x_test)

[72]: from sklearn.metrics import accuracy_score

[73]: accuracy_score(ypred,y_test)

[73]: 0.9559496567505721

[74]: userInput = np.array([[25,100,1.5,45]])

[75]: output = cls.predict(userInput)

[76]: output

[76]: array([0], dtype=int64)

[105]: if output[0] == 0:
      print("Failure is not predicted")
      else:
      print("Failure is predicted")

Failure is not predicted

[78]: # getting same and maximum accuracy by using linear and rbf kernal

[80]: # Gradient Boosting

[81]: from sklearn.datasets import load_iris
      from sklearn.model_selection import train_test_split
      from sklearn.ensemble import GradientBoostingClassifier

```

```

from sklearn.metrics import accuracy_score

# Split data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

# Initialize Gradient Boosting Classifier
gb_classifier = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1,
    random_state=42)

# Fit the model
gb_classifier.fit(X_train, y_train)

# Make predictions
y_pred = gb_classifier.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

```

Accuracy: 0.948512585812357

```
[ ]: # Experiment with ensemble techniques to improve model performance.
```

```

[101]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import VotingClassifier
from sklearn.metrics import accuracy_score

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

# Initialize individual classifiers
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)
gb_classifier = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1,
    random_state=42)
lr_classifier = LogisticRegression(max_iter=1000, random_state=42)

# Bagging (Random Forest)
bagging_classifier = RandomForestClassifier(n_estimators=100, random_state=42)

# Boosting (Gradient Boosting)
boosting_classifier = GradientBoostingClassifier(n_estimators=100,
    learning_rate=0.1, random_state=42)

```

```

# Stacking (Voting Classifier)
estimators = [('random_forest', rf_classifier), ('gradient_boosting',
    ↳gb_classifier), ('logistic_regression', lr_classifier)]
stacking_classifier = VotingClassifier(estimators)

# Train individual classifiers
rf_classifier.fit(X_train, y_train)
gb_classifier.fit(X_train, y_train)
lr_classifier.fit(X_train, y_train)

# Train ensemble classifiers
bagging_classifier.fit(X_train, y_train)
boosting_classifier.fit(X_train, y_train)
stacking_classifier.fit(X_train, y_train)

# Make predictions
y_pred_rf = rf_classifier.predict(X_test)
y_pred_gb = gb_classifier.predict(X_test)
y_pred_lr = lr_classifier.predict(X_test)
y_pred_bagging = bagging_classifier.predict(X_test)
y_pred_boosting = boosting_classifier.predict(X_test)
y_pred_stacking = stacking_classifier.predict(X_test)

# Evaluate individual classifiers
accuracy_rf = accuracy_score(y_test, y_pred_rf)
accuracy_gb = accuracy_score(y_test, y_pred_gb)
accuracy_lr = accuracy_score(y_test, y_pred_lr)

# Evaluate ensemble classifiers
accuracy_bagging = accuracy_score(y_test, y_pred_bagging)
accuracy_boosting = accuracy_score(y_test, y_pred_boosting)
accuracy_stacking = accuracy_score(y_test, y_pred_stacking)

print("Individual Classifier Accuracies:")
print("Random Forest:", accuracy_rf)
print("Gradient Boosting:", accuracy_gb)
print("Logistic Regression:", accuracy_lr)

print("\nEnsemble Classifier Accuracies:")
print("Bagging (Random Forest):", accuracy_bagging)
print("Boosting (Gradient Boosting):", accuracy_boosting)
print("Stacking (Voting Classifier):", accuracy_stacking)

```

```

Individual Classifier Accuracies:
Random Forest: 1.0
Gradient Boosting: 1.0
Logistic Regression: 1.0

```

Ensemble Classifier Accuracies:

Bagging (Random Forest): 1.0

Boosting (Gradient Boosting): 1.0

Stacking (Voting Classifier): 1.0

```
[100]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

# Initialize individual regressors
rf_regressor = RandomForestRegressor(n_estimators=100, random_state=42)
gb_regressor = GradientBoostingRegressor(n_estimators=100, learning_rate=0.1,
    random_state=42)
lr_regressor = LinearRegression()

# Train individual regressors
rf_regressor.fit(X_train, y_train)
gb_regressor.fit(X_train, y_train)
lr_regressor.fit(X_train, y_train)

# Make predictions
y_pred_rf = rf_regressor.predict(X_test)
y_pred_gb = gb_regressor.predict(X_test)
y_pred_lr = lr_regressor.predict(X_test)

# Calculate RMSE and MAE for each regressor
rmse_rf = mean_squared_error(y_test, y_pred_rf, squared=False)
mae_rf = mean_absolute_error(y_test, y_pred_rf)

rmse_gb = mean_squared_error(y_test, y_pred_gb, squared=False)
mae_gb = mean_absolute_error(y_test, y_pred_gb)

rmse_lr = mean_squared_error(y_test, y_pred_lr, squared=False)
mae_lr = mean_absolute_error(y_test, y_pred_lr)

print("Random Forest Regressor:")
print("RMSE:", rmse_rf)
print("MAE:", mae_rf)
print("\nGradient Boosting Regressor:")
print("RMSE:", rmse_gb)
print("MAE:", mae_gb)
```

```
print("\nLinear Regression:")
print("RMSE:", rmse_lr)
print("MAE:", mae_lr)
```

Random Forest Regressor:
RMSE: 0.037193189340702336
MAE: 0.013666666666666667

Gradient Boosting Regressor:
RMSE: 0.06593367853774076
MAE: 0.03101948917402032

Linear Regression:
RMSE: 0.1926494080135646
MAE: 0.1463769496530853

```
[ ]: # Compare the performance of different models and identify the best-performing
      ↪ one.
```

```
[102]: #from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

# Load the Iris dataset
#X, y = load_iris(return_X_y=True)

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪random_state=42)

# Initialize the classifiers
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)
gb_classifier = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1,
      ↪random_state=42)
lr_classifier = LogisticRegression(max_iter=1000, random_state=42)

# Train the classifiers
rf_classifier.fit(X_train, y_train)
gb_classifier.fit(X_train, y_train)
lr_classifier.fit(X_train, y_train)

# Make predictions
y_pred_rf = rf_classifier.predict(X_test)
y_pred_gb = gb_classifier.predict(X_test)
y_pred_lr = lr_classifier.predict(X_test)
```

```

# Calculate accuracy for each classifier
accuracy_rf = accuracy_score(y_test, y_pred_rf)
accuracy_gb = accuracy_score(y_test, y_pred_gb)
accuracy_lr = accuracy_score(y_test, y_pred_lr)

print("Accuracy Scores:")
print("Random Forest:", accuracy_rf)
print("Gradient Boosting:", accuracy_gb)
print("Logistic Regression:", accuracy_lr)

# Identify the best-performing model
best_model = max(accuracy_rf, accuracy_gb, accuracy_lr)
if best_model == accuracy_rf:
    print("Best Performing Model: Random Forest")
elif best_model == accuracy_gb:
    print("Best Performing Model: Gradient Boosting")
else:
    print("Best Performing Model: Logistic Regression")

```

Accuracy Scores:
 Random Forest: 1.0
 Gradient Boosting: 1.0
 Logistic Regression: 1.0
 Best Performing Model: Random Forest

```
[ ]: # Perform cross-validation to assess the model's generalization ability.
```

```

[104]: from sklearn.model_selection import cross_val_score
        from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
        from sklearn.svm import SVC

        # Initialize the classifiers
        rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)
        svm_classifier = SVC()
        gb_classifier = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1,
        ↪random_state=42)

        # cross-validation with 5 folds
        cv_scores_rf = cross_val_score(rf_classifier, X, y, cv=5)
        cv_scores_svm = cross_val_score(svm_classifier, X, y, cv=5)
        cv_scores_gb = cross_val_score(gb_classifier, X, y, cv=5)

        # Print the cross-validation scores
        print("Cross-Validation Scores for Random Forest Classifier:", cv_scores_rf)
        print("Mean Cross-Validation Score for Random Forest Classifier:", cv_scores_rf.
        ↪mean())

```

```

print("\nCross-Validation Scores for Support Vector Machine (SVM):",  

      ↪cv_scores_svm)  

print("Mean Cross-Validation Score for Support Vector Machine (SVM):",  

      ↪cv_scores_svm.mean())  

print("\nCross-Validation Scores for Gradient Boosting Classifier:",  

      ↪cv_scores_gb)  

print("Mean Cross-Validation Score for Gradient Boosting Classifier:",  

      ↪cv_scores_gb.mean())

```

Cross-Validation Scores for Random Forest Classifier: [0.96666667 0.96666667
0.93333333 0.96666667 1.]

Mean Cross-Validation Score for Random Forest Classifier: 0.9666666666666668

Cross-Validation Scores for Support Vector Machine (SVM): [0.96666667 0.96666667
0.96666667 0.93333333 1.]

Mean Cross-Validation Score for Support Vector Machine (SVM): 0.9666666666666666

Cross-Validation Scores for Gradient Boosting Classifier: [0.96666667 0.96666667
0.9 0.96666667 1.]

Mean Cross-Validation Score for Gradient Boosting Classifier: 0.9600000000000002

```
[ ]: # Cross Validation is performed
```